

Project 2 – Finding the Closest Pair

Qiang Huang

huangq25@mail2.sysu.edu.cn

School of Data and Computer Science


Sun Yat-Sen University

Introduction

- The purpose of the project:
 - Implement an **approximation method** with random projection to find the **closest pair** of a set of objects in high-dimensional space;
 - Use **Euclidean distance** as the distance metric.

Closest Pair Problem

- Given n objects $\{o_1, o_2, \dots, o_n\}$ in d -dimensional Euclidean space R^d , the **closest pair problem** is to find a pair of objects o_i and o_j that minimizes $\|o_i - o_j\|$.



The closest pair is $\{o_2, o_6\}$

Euclidean Distance

- Given two objects $\mathbf{o}_i = (o_{i1}, o_{i2}, \dots, o_{id})$ and $\mathbf{o}_j = (o_{j1}, o_{j2}, \dots, o_{jd})$ from R^d , the Euclidean distance between o_i and o_j is computed as follows:

$$\|\mathbf{o}_i - \mathbf{o}_j\| = \sqrt{\sum_{k=1}^d (o_{ik} - o_{jk})^2}$$

- An Example:
 - Suppose $o_i = (4, 2)$ and $o_j = (1, 6)$, then
 - $\|\mathbf{o}_i - \mathbf{o}_j\| = \sqrt{(4 - 1)^2 + (2 - 6)^2} = \sqrt{3^2 + 4^2} = 5$

Overview of the Algorithm

- Step 1: Make a **random projection** to project the n data objects from R^d to m random lines $\{S_1, S_2, \dots, S_m\}$;
- Step 2: For each random line S_i , we find the closest pair cp as a candidate; (**closest pair in $R^d \rightarrow$ closest pair in 1-Dimension**)
- Step 3: We compute the Euclidean distance for the m candidates (each candidate is the closest pair of S_i), and return the closest pair among the m candidates;

Overview of the Algorithm

Algorithm 1: Closest-Pair


Input: a database D of n data objects
 $\{o_1, o_2, \dots, o_n\}$ and the number of
random projections m .

Output: the closest pair cp .

```
1  $\{S_1, S_2, \dots, S_m\} = \text{Random-Projection()};$ 
2 Let  $cp$  be the closest pair,  $cp = \text{null}$ ;
3 Let  $min$  be the closest distance,  $min = \infty$ ;
4 for  $i = 1$  to  $m$  do
5      $\{\delta, cp_i\} = \text{Closest-Pair-Line}(S_i);$ 
6     Compute Euclidean distance  $dist$  for the
       closet pair  $cp_i$  found in  $S_i$ ;
7     if  $dist < min$  then
8          $min = dist$ ;
9          $cp = cp_i$ ;
10 return  $cp$ ;
```

Random Projection

- Project all n objects onto the m random lines S_1, S_2, \dots, S_m ;



The random projection is computed by $v_{ij} = \vec{a}_i \cdot \vec{o}_j$

Random Projection

- Generate random projection vectors $\vec{a}_1, \vec{a}_2, \dots, \vec{a}_m$
 - Each \vec{a}_i is a d -dimensional vector where each component is drawn from **standard Normal distribution $N(0,1)$** .
- There is a **Box-Muller method** to generate a random variable from $N(0,1)$.
 - Suppose U_1 and U_2 are random variable distributed uniformly from $(0,1)$;
 - Two independent random variables X_1 and X_2 from $N(0,1)$ are generated as follows:
 - $X_1 = \sqrt{-2 \log U_1} \cos(2\pi U_2)$
 - $X_2 = \sqrt{-2 \log U_1} \sin(2\pi U_2)$

Random Projection


Algorithm 2: Random-Projection

Input: a database D of n data objects
 $\{o_1, o_2, \dots, o_n\}$ and the number of
random projections m .

Output: $\{S_1, S_2, \dots, S_m\}$.


- 1 Generate a set of m random projection vectors
 $\{\vec{a}_1, \vec{a}_2, \dots, \vec{a}_m\}$;
 - 2 $S_j = \emptyset$ for $j \in \{1, 2, \dots, m\}$;
 - 3 **for** $i = 1$ *to* n **do**
 - 4 **for** $j = 1$ *to* m **do**
 - 5 $c_{i,j} = i$;
 - 6 $v_{i,j} = \vec{a}_j \cdot \vec{o}_i$;
 - 7 $S_j = S_j \cup (c_{i,j}, v_{i,j})$;
 - 8 **return** $\{S_1, S_2, \dots, S_m\}$;
-

Closest Pair in 1-Dimension




- 1D problem can be solved in $O(n \log n)$ via sorting (i.e. quick sort).
- We can sort all objects on the line, and find the closest pair with a linear scan.
 - i.e., the closest pair is $\{o_3, o_5\}$ for the random line \vec{a}_1
- However, can we find the closest pair during the sorting?

1-Dimension Divide and Conquer



- Yes, use divide-and-conquer!
- Divide:
 - Divide all the objects into two sets S_1 and S_2 by a **median m** so that $p < q$ for all $p \in S_1$ and $q \in S_2$;
 - Recursively compute the closest pair $\{p_1, p_2\}$ in S_1 and $\{q_1, q_2\}$ in S_2 .

1-Dimension Divide and Conquer



- Conquer:

- Let δ be the smallest separation found so far:

$$\delta = \min(|p_2 - p_1|, |q_2 - q_1|)$$

- The closest pair is $\{p_1, p_2\}$, or $\{q_1, q_2\}$, or $\{p_3, q_3\}$.
- In 1-Dimension, p_3 is the rightmost object of S_1 and q_3 is the leftmost object of S_2 (Why?)

1-Dimension Divide and Conquer

Algorithm 3: Closest-Pair-Median(S)

Input: a set S of n data objects on a line

Output: the smallest separation δ and the closest pair cp

```
1 if  $|S| = 1$  then
2   return  $\{\delta = \infty, cp = null\}$ ;
3 if  $|S| = 2$  then
4   if  $p_1 > p_2$  then swap( $p_1, p_2$ );
5   return  $\{\delta = |p_2 - p_1|, cp = \{p_1, p_2\}\}$ ;
6 Let  $m = median(S)$ ;
7 Divide  $S$  into  $S_1$  and  $S_2$  at  $m$ ;
8  $\{\delta_1, cp_1\} = \text{Closest-Pair-Median}(S_1)$ ;
9  $\{\delta_2, cp_2\} = \text{Closest-Pair-Median}(S_2)$ ;
10  $\{\delta_{12}, cp_{12}\}$  is the result across the cut;
11 return  $\{\delta, cp\} =$ 
     $\min(\{\delta_1, cp_1\}, \{\delta_2, cp_2\}, \{\delta_{12}, cp_{12}\})$ ;
```

Finding the Median

- We can use the divide-and-conquer to find the median in $O(n)$.
- Given a set of objects S
- **Select(S, k):**
 - choose a splitter $a_i \in S$ **uniformly at random**; (add randomization)
 - Split S into two sub sets S^- and S^+ ; (divide)
 - For all $a_j \in S^-$, $a_j < a_i$;
 - For all $a_j \in S^+$, $a_j > a_i$;
 - If $|S^-| = k - 1$ then a_i is the answer; (conquer)
 - Else if $|S^-| \geq k$ then **Select(S^-, k)**;
 - Else **Select($S^+, k - 1 - |S^-|$)**;

Finding the Median


- An Example:
- Suppose $S = \{4, 1, 7, 3, 5, 2, 6\}$ and $k = 6$.
- Find a splitter $a_i = 3$, then $S = \{2, 1, 3, 4, 5, 7, 6\}$,
- Since $|S^-| = 2 < k - 1 = 5$, consider $S^+ = \{4, 5, 7, 6\}$ and $\underline{k = 6 - 1 - |S^-| = 3}$
- Find a splitter $a_i = 6$, then $S = \{2, 1, 3, 4, 5, 6, 7\}$,
- Since $|S^-| = 2 = k - 1$, a_i is the answer

Finding the Median

- Why the algorithm $\text{Select}(S, k)$ is $O(n)$? (average case, not worst case)
- choose a splitter $a_i \in S$ ***uniformly at random*** (add a **randomization**)
- In each recursion j (Event X_j), we can prune S **from $n(\frac{3}{4})^{j-1}$ to $n(\frac{3}{4})^j$** (the best case is $\frac{1}{2}$, here $\frac{3}{4}$ is the average case)
 - i.e., $n: 7 \rightarrow 4$ after 1st recursion in the example of the previous slide.
- When j increases, $n(\frac{3}{4})^j \rightarrow 1$; Finally we can find the k^{th} object.
- Thus, the average time is
- $E[x] = \sum_j E[X_j] \leq \sum_j cn(\frac{3}{4})^j = cn \sum_j (\frac{3}{4})^j = O(n)$ (here $\sum_j (\frac{3}{4})^j = 4$)

1-Dimension Divide and Conquer

- Shall we have to choose **median** for splitting?
- Actually we can generalize the **median** to a **random pivot**:
 - choose a pivot $p_i \in S$ uniformly at random;
 - Split S into two sub sets S_1 and S_2 ;
 - For all $p_j \in S_1, p_j \leq p_i$;
 - For all $p_j \in S_2, p_j > p_i$;



1-Dimension Divide and Conquer

Algorithm 4: Closest-Pair-Pivot(S)

Input: a set S of n data objects on a line

Output: the smallest separation δ and the closest pair cp

```
1 if  $|S| = 1$  then
2   return  $\{\delta = \infty, cp = null\}$ ;
3 if  $|S| = 2$  then
4   if  $p_1 > p_2$  then swap( $p_1, p_2$ );
5   return  $\{\delta = |p_2 - p_1|, cp = \{p_1, p_2\}\}$ ;
6 Choose a pivot  $p_i \in S$  uniformly at random;
7 for each element  $p_j \in S$  do
8   Put  $p_j$  in  $S_1$  if  $p_j \leq p_i$ ;
9   Put  $p_j$  in  $S_2$  if  $p_j > p_i$ ;
10  $\{\delta_1, cp_1\} = \text{Closest-Pair-Pivot}(S_1)$ ;
11  $\{\delta_2, cp_2\} = \text{Closest-Pair-Pivot}(S_2)$ ;
12  $\{\delta_{12}, cp_{12}\}$  is the result across the cut;
13 return  $\{\delta, cp\} =$ 
    min( $\{\delta_1, cp_1\}, \{\delta_2, cp_2\}, \{\delta_{12}, cp_{12}\}$ );
```

Dataset

- Please download from <http://yann.lecun.com/exdb/mnist/>
- Use TRAINING SET IMAGE FILE (train-images-idx3-ubyte) as dataset Mnist.ds ($n = 60,000$ and $d = 784$)
- Each object is an image with pixels 28×28
- Please read the file format on the website and extract the dataset from the train-images-idx3-ubyte.gz

Input Format of the Project


- The input format of the dataset is described as follows:

```
1 object1,1 object1,2 ... object1,d
2 object2,1 object2,2 ... object2,d
...
N objectN,1 objectN,2 ... objectN,d
```

- The 1st element is the object id, the 2nd element to the $(d + 1)^{\text{th}}$ element are the coordinate of the object itself.

Requirements

- Implementations:
 - We limit $m = 100$
 - Programming language: C/C++
 - implement **Algorithm 1** with **Algorithm 2&3**, and display the running time and the closest pair
 - implement **Algorithm 1** with **Algorithm 2&4**, and display the running time and the closest pair
 - the **running time** is defined by the wall clock time of Algorithm 1 to find the closest pair.
 - the closest pair can be displayed as the right figure:
- Write an experimental report:
 - Depict how to develop the two implementations
 - Discuss these two implementations with the running time and your analyses.



Submission

- 提交时间：2016年12月11日 23:59分
- 提交方式：提交给各方向学委，学委统一交给TA（张楚涵）
- 提交文件：
 - **README**
 - **源代码**: 保存在目录 `“/src”`
 - **数据集**: 假设保存在 `“/src”`（不用提交数据集）
 - **Makefile**: 如果你使用了多个.h 和.cpp文件，需要写一个makefile来编译，并生成一个可执行的程序 `cp`。程序必须使用以下命令执行：

```
./cp -n 60000 -d 784 -f Mnist.ds
```
 - 实验报告（列明小组成员及分工）
 - **请注意不要抄袭！**
- 所有文件打包成一个.zip文件（不接受.rar）
 - 命名规范：“小组序号_组长学号_组长姓名（拼音）_Project2.zip”。
 - i.e., 1_14141414_zhangsan_Project2.zip